Kanchan Chowdhury

Email : kchowdh1@asu.edu
Webpage: kanchanchy.github.io
Mobile : +1-480-410-8677
Linkedin: linkedin.com/in/kanchan-chowdhury-5729699a

Research Interests

Machine Learning Systems, Database Systems, and Geospatial Data Analytics

Education

- **Arizona State University** — Tempe, Arizona
  PhD & MS in Computer Science — *Aug. 2018 – July 2024*
  Advisor: Prof. Mohamed Sarwat & Prof. Jia Zou, CGPA: 4.00

- **Chittagong University of Engineering and Technology** — Chittagong, Bangladesh
  Bachelor of Science in Computer Science and Engineering — *Mar. 2010 – Nov. 2014*
  Advisor: Prof. Mohammed Moshiul Hoque, CGPA: 3.76

Experience

- **Marquette University** — Milwaukee, Wisconsin
  Assistant Professor — *Aug. 2025 - Present*
  Department of Computer Science

- **Arizona State University** — Tempe, Arizona
  Postdoctoral Research Scholar — *Aug. 2024 - Jun. 2025*
  Mentor: Prof. Jia Zou
  - **Responsibilities**:
    1) Conducting research on optimizing end-to-end pipelines consisting of database queries and machine learning models for inference workloads, reducing runtime and memory usage of in-database machine learning systems.
    2) Helping my mentor in writing research proposals for grants, reviewing research papers, and co-mentoring students.

- **Arizona State University** — Tempe, Arizona
  Research Assistant — *Aug. 2018 - July 2024*
  - **Research Projects**:
    1) Co-optimization of machine learning and join queries based on model decomposition and join push-down.
    2) Designing and implementing a deep-learning and data processing system for raster imagery and vector datasets.
    3) Re-partitioning geospatial datasets to reduce spatial model training time and memory usage.

- **Wherobots Inc.** — Scottsdale, Arizona
  Research and Development Intern — *Jan. 2023 - Aug. 2023*
  - **Responsibilities**: Designing and developing spatial machine learning and deep learning tools, scalable map-matching, and geospatial data analytical algorithms. Integrating the developed tools into the Wherobots cloud platform.

- **Gagagugu PTE LTD** — Dhaka, Bangladesh
  Software Engineer — *Jan. 2017 - Jun. 2018*
  - **Responsibilities**: Developing Android Apps with social networking features such as calling, messaging, and posting.

- **Le Chef Plc** — Dhaka, Bangladesh
  Android Application Developer — *Jan. 2015 - Dec. 2016*
  - **Responsibilities**: Developing Android Apps featuring online order and reservation services for restaurants in the UK.

Teaching

- **Instructor** Fall 2022
  Spatial Data Science and Engineering (CSE 594) *Arizona State University*
  - **Responsibilities**:
    1) Preparing lecture slides and conducting lectures.
    2) Preparing assignments, projects, exam questions, and grading rubrics.
    3) Office hours to help students understand lectures and projects.
- **Teaching Assistant** Fall 2021, Spring 2021, Fall 2020, and Spring 2020
  Distributed Database Systems (CSE 512) *Arizona State University*
- **Teaching Assistant** Fall 2019
  Data Processing at Scale (CSE 511) *Arizona State University*
- **Teaching Assistant** Spring 2019
  Object-Oriented Programming & Data Structure (CSE 205) *Arizona State University*
- **Teaching Assistant** Spring 2019
  Principles of Programming with C++ (CSE 100) *Arizona State University*
- **Teaching Assistant** Fall 2018
  Principles of Programming with Java & Python (CSE 100) *Arizona State University*

Publications

- **Kanchan Chowdhury**, Lulu Xie, Lixi Zhou, Jia Zou; ExBoost: Out-of-Box Co-Optimization of Machine Learning and Join Queries. *To Appear in DASFAA 2025*
- Hong Guan, Lixi Zhou, Lei Yu, Li Xiong, **Kanchan Chowdhury**, Lulu Xie, Xusheng Xiao, Jia Zou; Privacy and Accuracy-Aware AI/ML Model Deduplication. *To Appear in ACM SIGMOD 2025*
- Hong Guan, Ansh Tiwari, Summer Gautier, Rajan Hari Ambrish, Lixi Zhou, Deepti Gupta, Yancheng Wang, Yingzhen Yang, Chaowei Xiao,, **Kanchan Chowdhury**, Jia Zou; Declarative Privacy-Preserving Inference Queries. *To Appear in DASFAA 2025 (Demo Track)*
- **Kanchan Chowdhury**, Mohamed Sarwat; Deep Learning with Spatiotemporal Data: A Deep Dive into GeotorchAI. *40th International Conference on Data Engineering (ICDE)*, 2024
- Lixi Zhou, Qi Lin, **Kanchan Chowdhury**, Saif Masood, Jia Zou; Serving Deep Learning Models from Relational Databases. *27th International Conference on Extending Database Technology (EDBT)*, 2024
- **Kanchan Chowdhury**, Mohamed Sarwat; A Demonstration of GeoTorchAI: A Spatiotemporal Deep Learning Framework. *ACM SIGMOD International Conference on Management of Data*, 2023 (Demo Track)
- **Kanchan Chowdhury**, Vamsi Meduri, Mohamed Sarwat; A Machine Learning-Aware Data Re-partitioning Framework for Spatial Datasets. *38th International Conference on Data Engineering (ICDE)*, 2022
- **Kanchan Chowdhury**, Mohamed Sarwat; GeoTorch: A Spatiotemporal Deep Learning Framework. *30th International Conference on Advances in Geographic Information Systems (SIGSPATIAL)*, 2022 (Short Paper)
- Vamsi Meduri, **Kanchan Chowdhury**, Mohamed Sarwat; Evaluation of Machine Learning Algorithms in Predicting the Next SQL Query From the Future. *ACM Transactions on Database Systems (TODS)*, 2021
- Jia Yu, **Kanchan Chowdhury**, Mohamed Sarwat; Tabula in Action: A Sampling Middleware for Interactive Geospatial Visualization dashboards. *Proceedings of the VLDB Endowment 13*, 2020 (Demo Track)
- Vamsi Meduri, **Kanchan Chowdhury**, Mohamed Sarwat; Recurrent Neural Networks for Dynamic User Intent Prediction in Human-Database Interaction. *International Conference on EDBT, 2019 (Short Paper)*
- **Kanchan Chowdhury**, Lamia Alam, Shyla Sarmin, Safayet Arefin, Mohammed Moshiul Hoque; A Fuzzy Features Based Online Handwritten Bangla Word Recognition Framework. *18th ICCIT*, 2015

Publications Under Review

- **Kanchan Chowdhury**, Lixi Zhou, Jia Zou; InferF: Declarative Factorization of AI/ML Inferences over Joins. *Under review of an A* data management conference.*
- Lixi Zhou, **Kanchan Chowdhury**, Jia Zou; CactusDB: Unlock Co-Optimization Opportunities for SQL Queries Nested with AI/ML Models. *Under review of an A* data management conference.*
- Lulu Xie, **Kanchan Chowdhury**, Jia Zou; A Model-Guided Framework for Generating Synthetic Identity Documents to Evaluate Fraud Detection Tools. *Under review of an A* machine learning conference.*

Reviewing Services

- Got invited and accepted to serve in the review board of VLDB 2026 starting from October 2025.
- **Conference Paper Reviewer**: Served as a reviewer for the conferences ICCAD 2023 and DAC 2025.
- **Journal Paper Reviewer**: Served as a reviewer for the journal IEEE TKDE, Scientific Data, and Journal of Advances in Information Technology.
- **External Reviewer**: Reviewed papers as an external reviewer for the following conferences and journals - SIGMOD 2020, 2022, 2025, PVLDB 2019-2025, ICDE 2020, SIGSPATIAL 2021, VLDB Journal, and TSAS Journal.

Additional Contributions

- **Grant Reviewer**: Served as a travel grant reviewer for the Graduate and Professional Student Association (GPSA) at Arizona State University from May 2022 to August 2023.
- **Conference Volunteer**: Volunteered to organize two conferences - SIGSPATIAL 2022 and SIGMOD 2023.
- **Presentations & Talks**: Five conference presentations - SIGMOD 2023, FOSS4GNA 2023, ICDE 2022, SIGSPATIAL 2022, and ICCIT 2015.
- **Open Source Contribution**: Contributed to Apache Sedona, an open-source geospatial cluster computing framework with 2k+ GitHub Stars, by adding support for two new spatial data types.
- **Membership**: I hold a membership in the professional organizations such as ACM, IEEE, and SIGMOD

Participation and Awards

- Recipient of ACM SIGMOD 2023 student travel award to attend the conference and present a paper.
- Recipient of ACM SIGSPATIAL 2022 travel award to attend the conference and present a paper.
- Received ASU Graduate and Professional Student Association (GPSA) travel grant twice - 2022 and 2023
- Recipient of CIDSE Doctoral Fellowship at Arizona State University for the academic year 2018-2019.
- 2nd Runner-up at National Hackathon organized by ICT Division of Bangladesh in 2014. The challenge of the hackathon was to design a project-based solution to solve a national problem of the country.
- 2nd Runner-up at National Mobile Application Code Hub organized by BUET, Bangladesh in 2014.
- Recipient of Honors award from my undergraduate university for maintaining academic excellence.
- Recepient of merit scholarship in all four years of my undergraduate education.
- Recepient of the government merit scholarship in Secondary School Certificate examination, Higher Secondary School Certificate examination, 8th-grade public examination, and 5th-grade public examination.

Technical Skills

- **Programming**: Python, Java, C, C++, Scala, SQL, and HTML
- **Databases**: PostgreSQL, SparkSQL, and MySQL
- **Machine Learning**: PyTorch, Scikit-learn, Keras, TensorFlow, ML & DL Models, and ML Statistics
- **Data Analytics**: Apache Spark, Apache Sedona, PySpark, GeoPandas, Pandas, Matplotlib, and Plotly
- **Teaching**: Course Curriculum and Course Material Development, Managing Courses on Canvas and D2L
- **Others**: Distributed Computing, SDE Design Principles, CI/CD Pipeline, Docker, Jira, and Rest API

Projects

- **GeoTorchAI**: A deep learning and data preprocessing framework for raster imagery and spatiotemporal vector datasets, with **400+ GitHub Stars**. It enables spatiotemporal machine learning practitioners to easily and efficiently implement spatiotemporal deep learning models, besides supporting scalable data preprocessing.
- **InferF**: An out-of-box AI/ML-SQL co-optimization approach for end-to-end inference workflows where the users specify a SQL query and a pre-trained model exported in ONNX format, and the end-to-end processing will be automatically optimized reducing the execution latency.
- **ML-Aware Data Re-partitioning**: This is a framework which aims at reducing the training time and memory usage of a spatial machine learning model by reducing the number of partitions in a spatial grid dataset. Experiments on four datasets achieved a significant reduction in training time and memory usage.
- **NLIDB-Bench**: A benchmark for evaluating state-of-the-art approaches of SQL query generation from natural language queries. It proposes a set of evaluation metrics and conducts experiments with four datasets.